

51.505 – Foundations of Cybersecurity

Week 12 – PKI

Created by **Pawel Szalachowski (2017)**

Modified by **Jianying Zhou (2018)**


Last updated: 30 Nov 2018

Recap

- Questions on Week 11's exercises?


Symmetric Keys (scalability)

- Scalability issues with symmetric crypto
 - ✓ Distribution
 - ✓ Challenges in managing n secrets



Public Key Infrastructure (PKI)

- Asymmetric crypto (DH, RSA, ...) solves the scalability problems, ... but creates a new one:
 - ✓ **How to ensure that public-key is accessible and authentic?**



Public-Key Infrastructure (PKI)

- An infrastructure that allows to recognize which public key belongs to whom
- There is a central authority, called the *Certificate Authority* (CA).
 - ✓ Everyone trusts the CA and knows its public key.
- Alice to join the PKI
 - ✓ Alice generates public/private key pair (PK_A, PK_A^-) and contacts the CA.
 - ✓ The CA verifies her identity and issues a signed *certificate* that claims that " PK_A belongs to Alice".
 - ✓ Alice can now contact Bob sending PK_A and the certificate.
 - ✓ As Bob trusts the CA, he trusts the certificate.

Trust and Trust Models

- Needed to solve scalability issues
- Trust Models
 - ✓ Centralized (Monopoly, Oligarchy)
 - ✓ Decentralized (Anarchy)

Trust and Trust Models

- **Monopoly**
 - ✓ Everyone trust a single organization to be the only CA in the world.
 - ✓ The key of that CA is embedded in all software and hardware as the PKI trust anchor.
- Problem:
 - ✓ There is no universally trusted entity.
 - Would a bank trust an external CA to issue certificates to its customers?

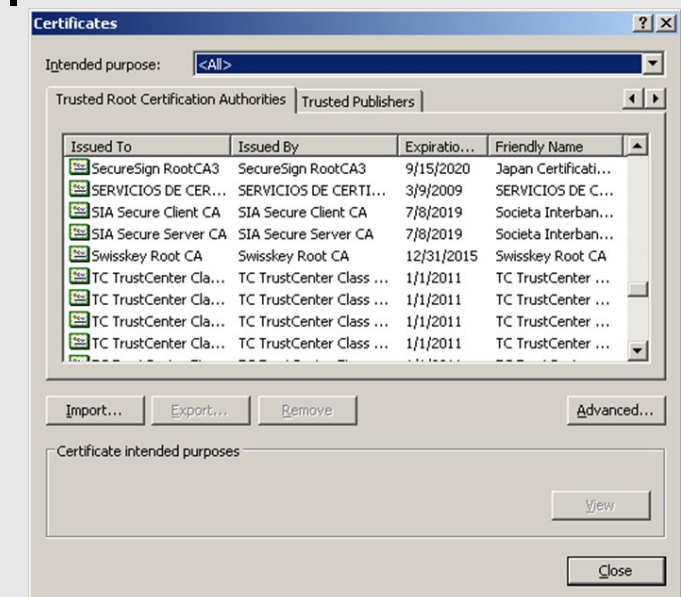
Trust and Trust Models

- **Oligarchy**

- ✓ Used in browsers.
- ✓ A number of organizations set themselves up as trust anchors.
 - Chances that one of these is corrupted is higher than the chance that a single trust anchor is corrupted.
- ✓ Product vendor selects amongst them.

- **Problem:**

- ✓ Make it easier to dupe users into accepting a trust anchor.
 - Even knowledgeable users will have a hard time checking the 80+ trust anchors in a current browser.



Trust and Trust Models

- **Anarchy**
 - ✓ Used in PGP.
 - ✓ Each user is responsible for configuring some trust anchors themselves.
 - Meeting people who hand out their certificates.
 - Search through public databases to find a path from one trust anchor to the key you want. (You implicitly trust everyone on that path!)
- **Problem:**
 - ✓ Not scalable, especially hard on certificate revocation/validation.
 - ✓ Mainly for personal (rather than corporate) use.

Trust and Trust Models


- Trust Models
 - ✓ Centralized (Monopoly, Oligarchy)
 - ✓ Decentralized (Anarchy)
- Impossible to have a universal PKI with a single trusted CA.
- The basic trust relationships are all based on contractual relationships.

PKI Examples

- SSL/TLS
 - ✓ Web (HTTPS), VPN, email, ...
- Credit card organizations
- Enterprises, companies, organizations, ...

Operations

- Registration Authority (RA)
 - ✓ verifies identities
- Certificate Authority (CA)
 - ✓ issues certificates
- Validation Authority (VA)
 - ✓ informs if a certificate is valid



Certificate

- Encoding of a particular data structure **must** be unique
 - ✓ X.509
- Fields
 - ✓ Subject: owner of the certificate
 - ✓ Issuer: issuer (CA) of the certificate
 - ✓ Not Before: the earliest time on which the certificate is valid
 - ✓ Not After: the latest time on which the certificate is valid
 - ✓ Public key: public key of the subject
 - ✓ Signature: signature of the certificate by the issuer's private key
- Other fields like serial number, key usage, algorithm id, ...
- Certificate can be extended for direct authorization.
 - ✓ Useful in a hierarchical CA structure, to limit the power of sub-CAs.

Multilevel Certificates

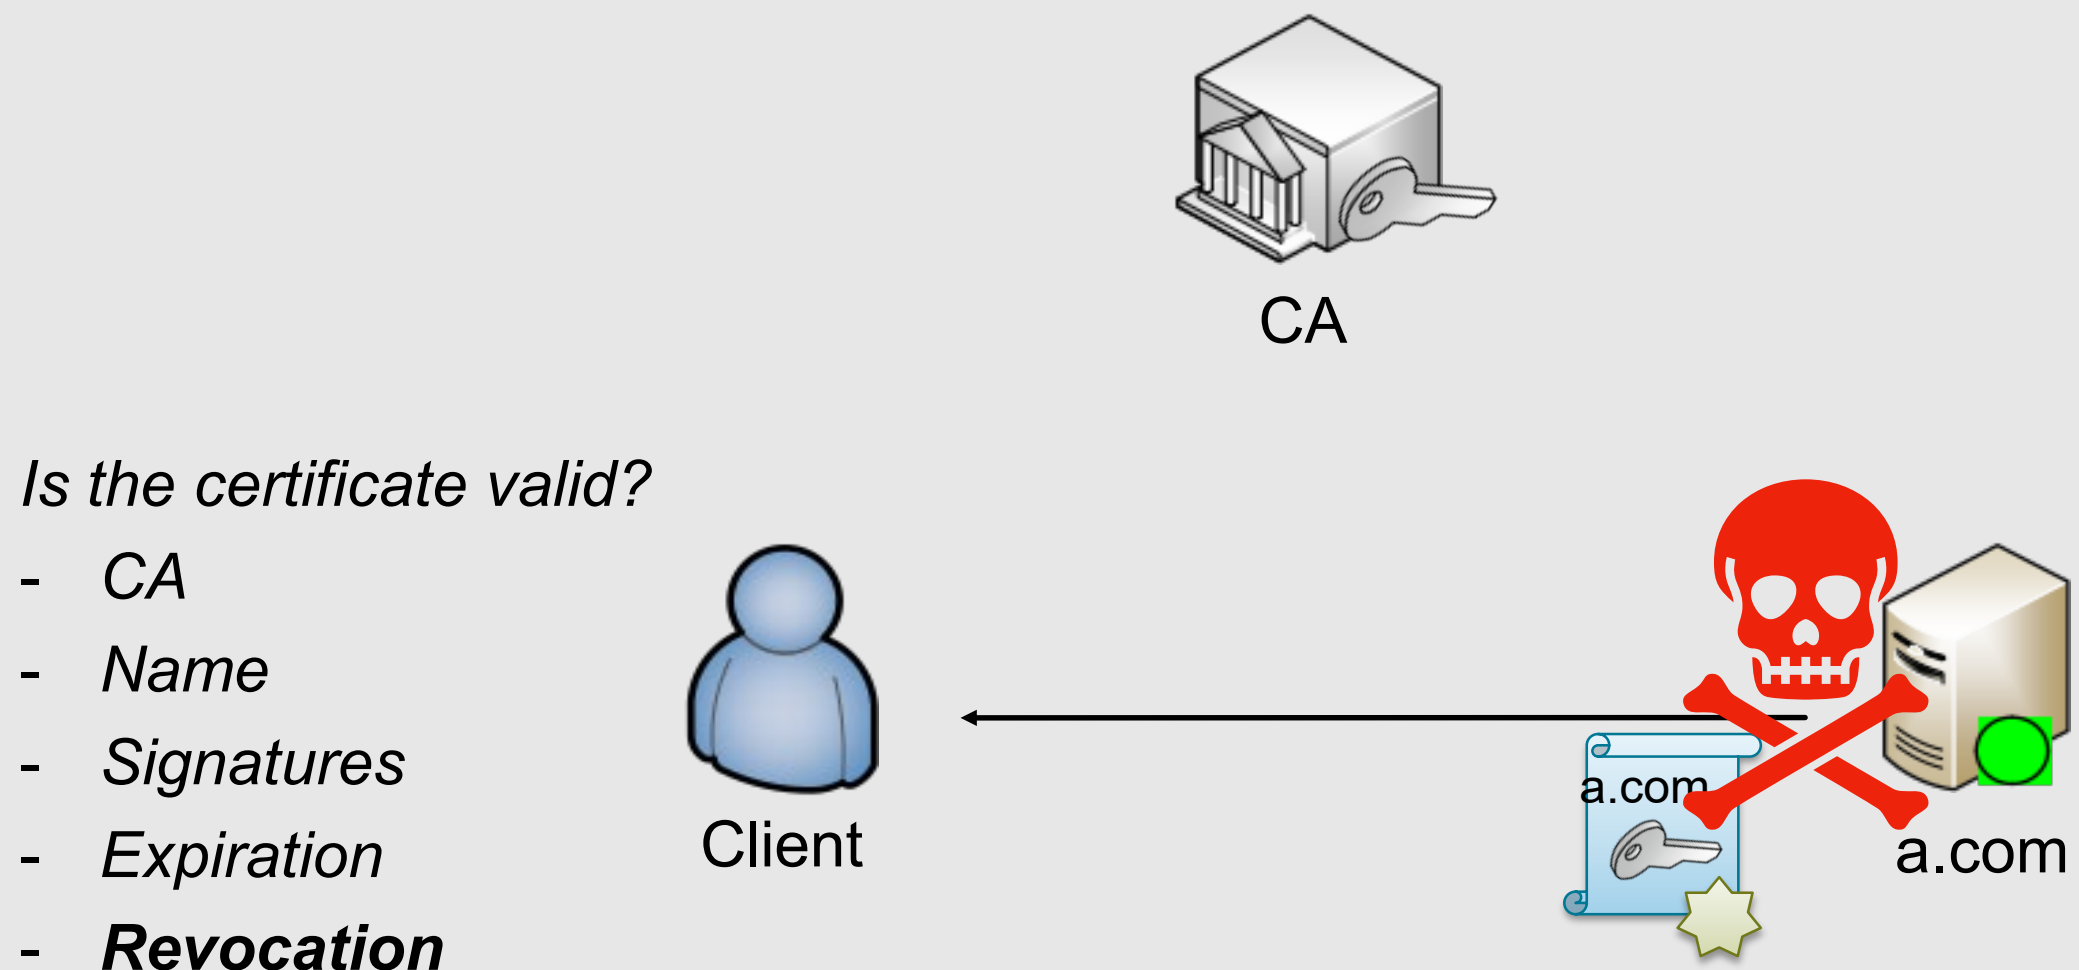
- For operational reasons certificates form chains
 - ✓ Root certificate (trust anchor)
 - self-signed certificate used for signing other certificates
 - ✓ Intermediate certificate
 - not self-signed, used for signing other certificates
 - ✓ Leaf certificate
 - cannot be used for signing other certificates

Certificate Revocation

- Sometimes a certificate has to be invalidated (revoked) by the issuing CA.
 - ✓ Reasons for certificate revocation?
 - ✓ How to do this? (One of the hardest problems to solve in a PKI.)
 - ✓ What if root/intermediate certificate has to be revoked? (Collateral damage.)
- Requirements:
 - ✓ Speed of revocation
 - ✓ Reliability of revocations
 - ✓ Overheads
 - ✓ Connectivity
- Possible solutions: revocation list, online verification, fast expiration

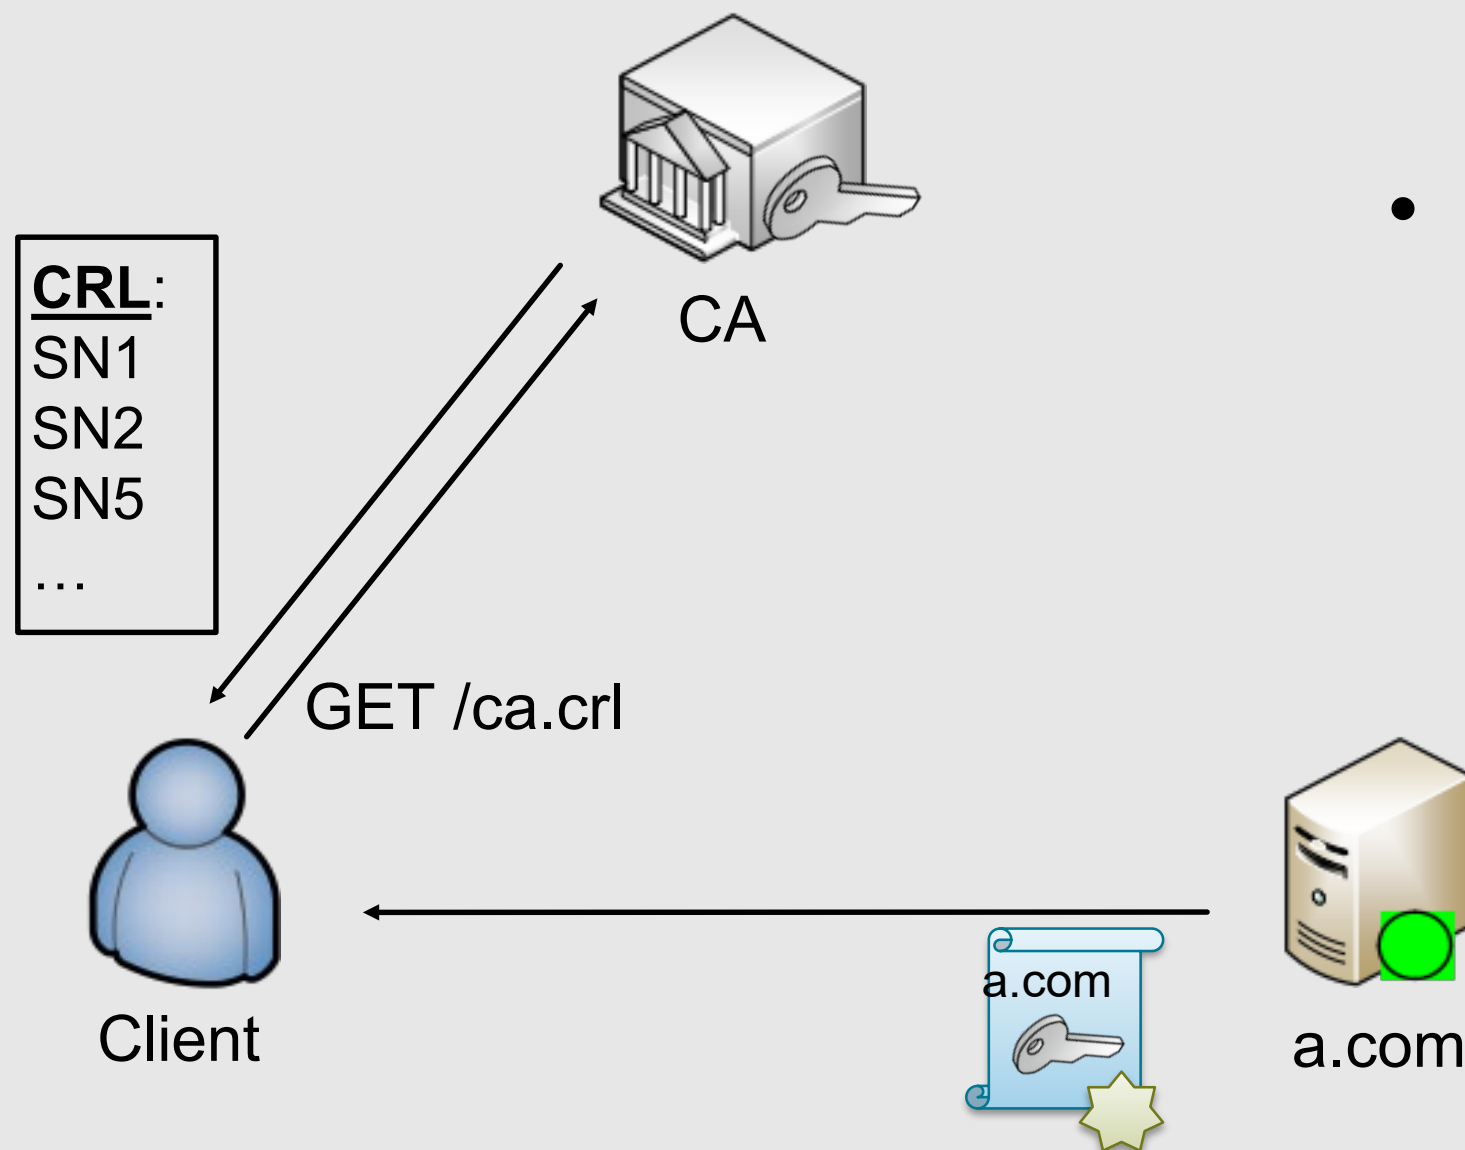
Current SSL/TLS PKI Model

- CA knows whether the certificate is valid or not.



Certificate Revocation List (CRL)

- CRL is a signed list of serial numbers that uniquely identify certificates being revoked.

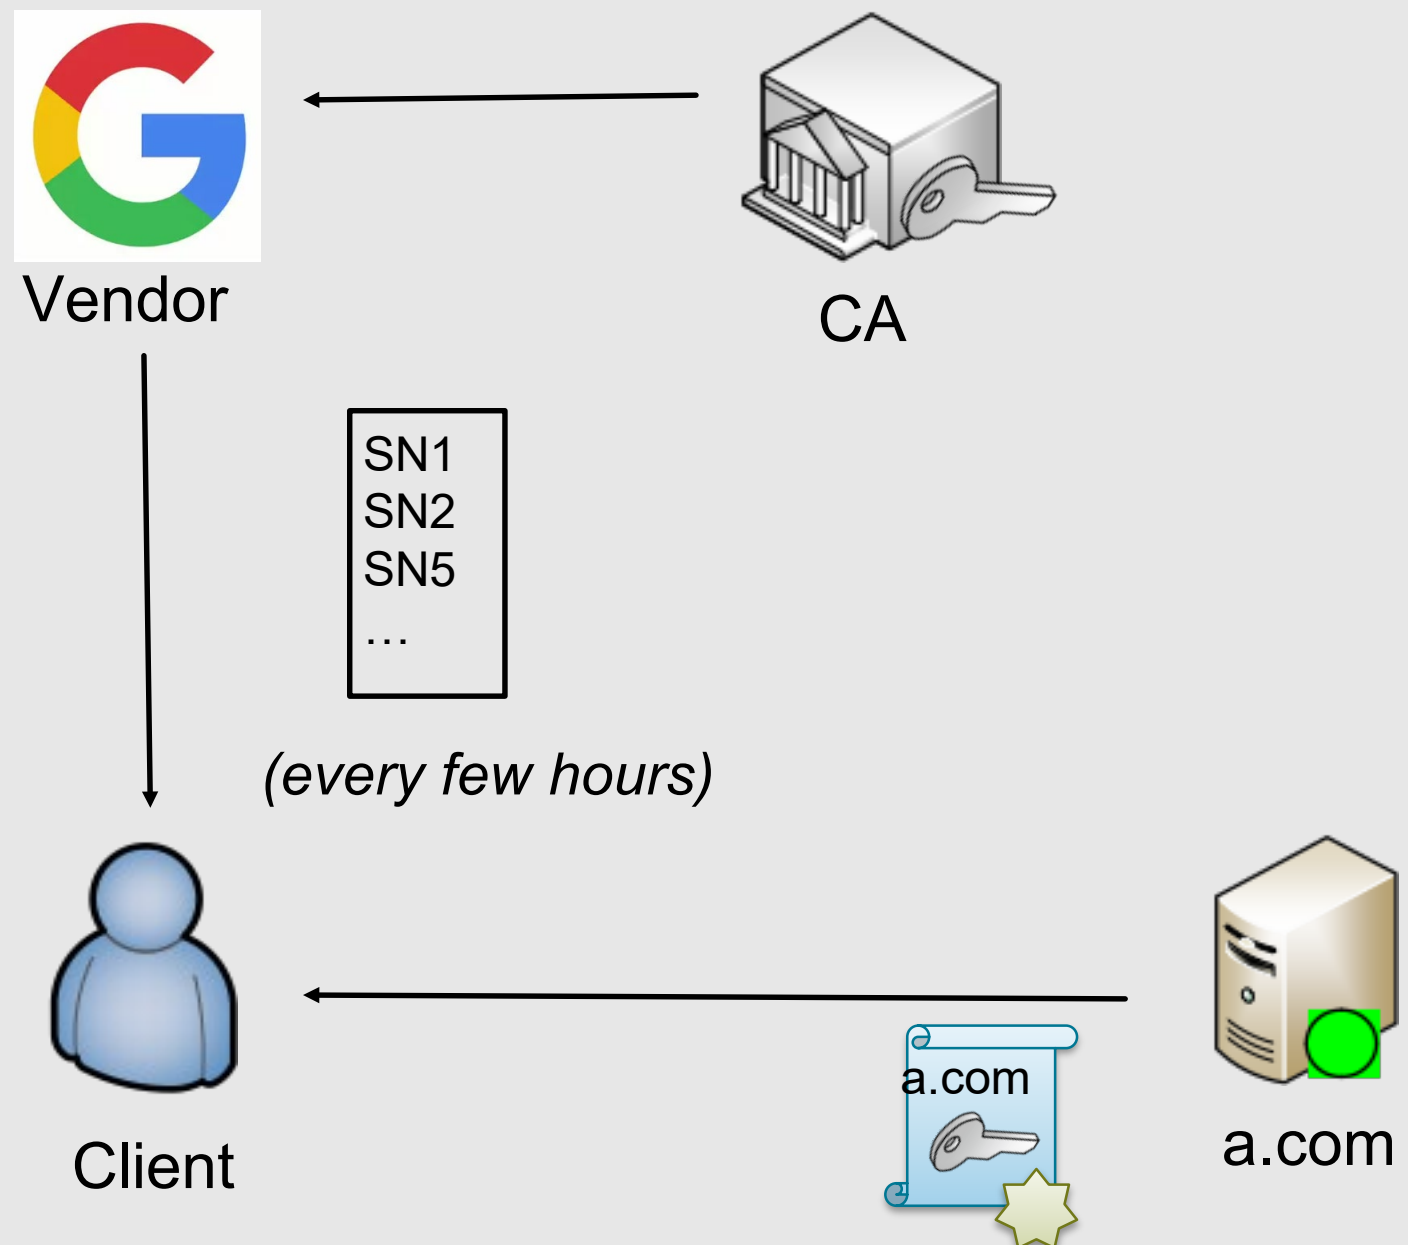


- **Problems:**

- ✓ CRLs are big (expensive)
- ✓ Slow revocation after key compromise/loss

CRLSets

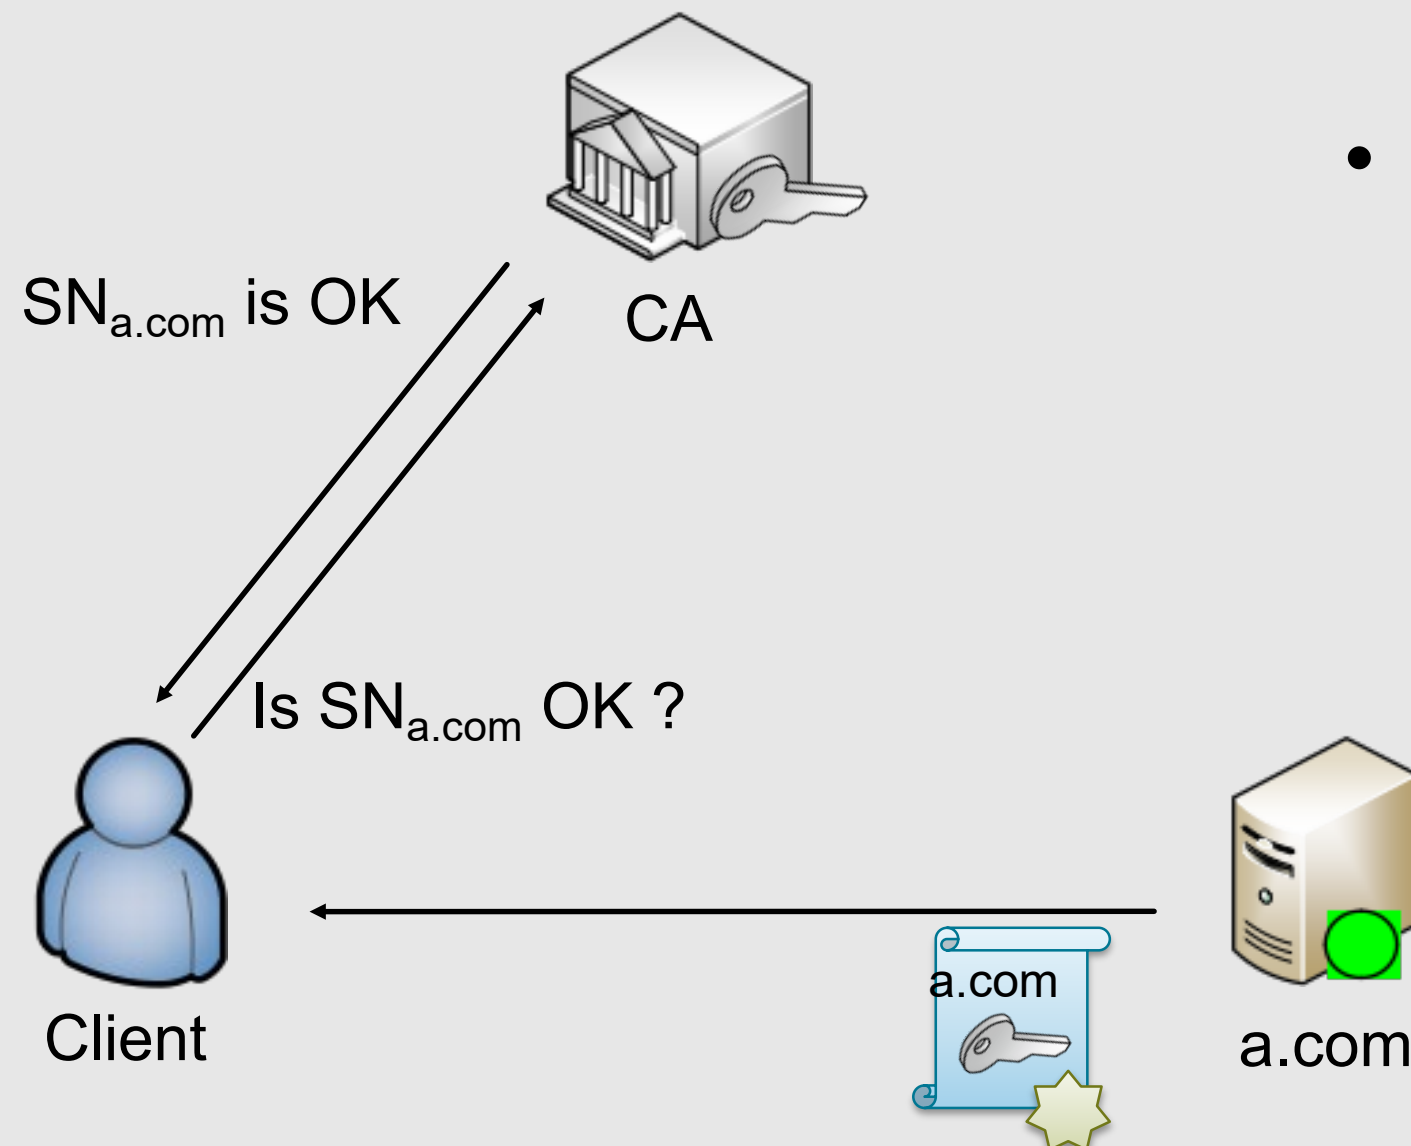
- CRLSets contains a carefully selected collection of revoked-certificate serial numbers published by many different certificate authorities.



- **Problems:**

- ✓ CRLSet is max 256KB
- ✓ Only 0.35% of all revocations are included
- ✓ Slow revocation after key compromise/loss

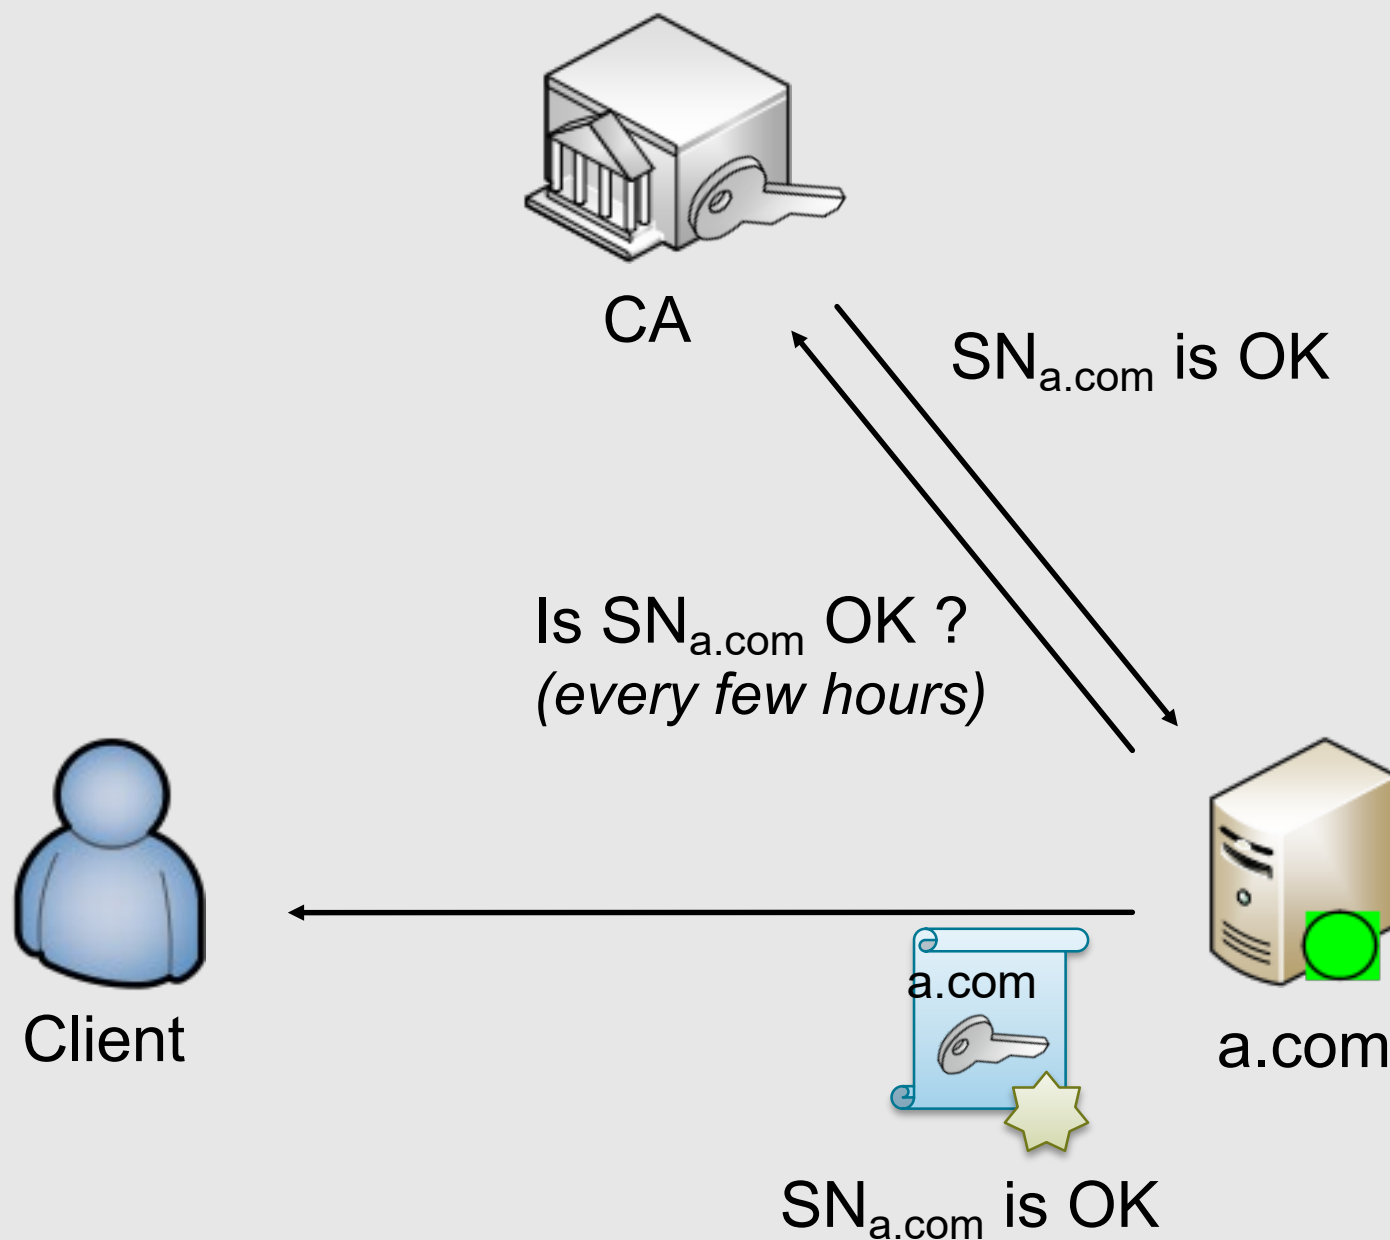
Online Certificate Status Protocol (OCSP)



- **Problems:**

- ✓ Blocking connection (~350ms)
- ✓ Availability (18% timeouts)

OCSP Stapling



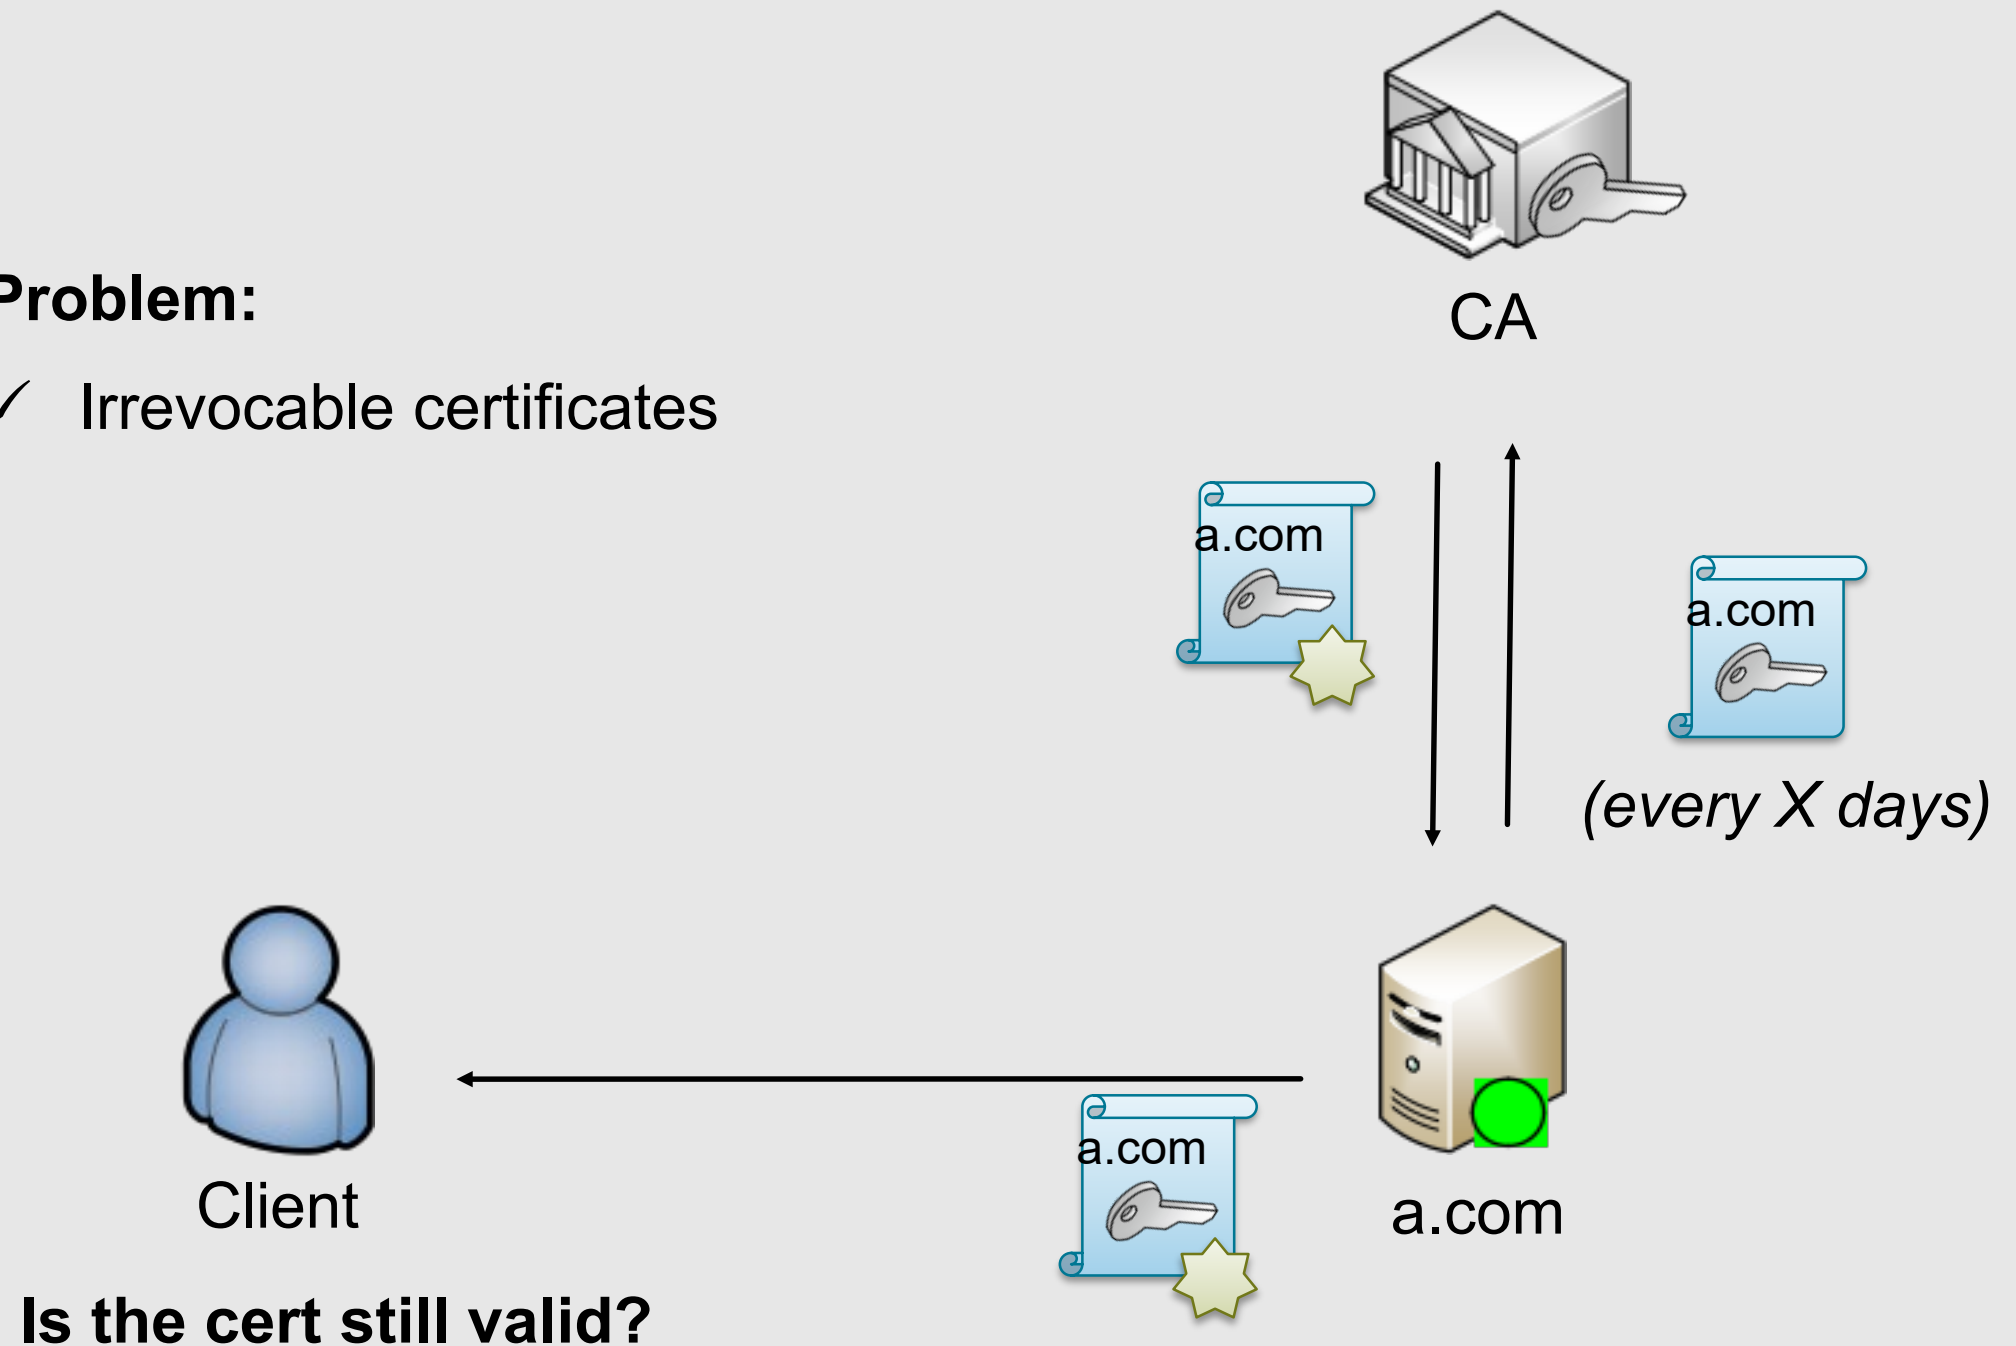
- **Problems:**

- ✓ Minimal server deployment < 3%
- ✓ Slow revocation after key compromise/loss

Short-lived Certificates

- **Problem:**

- ✓ Irrevocable certificates



Current State

		Desktop Browsers									Mobile Browsers				
		Chrome 44			Firefox	Opera		Safari	IE		iOS	Andr. 4.1–5.1		IE	
		OS X	Win.	Lin.	40	12.17	31.0	6–8	7–9	10	11	6–8	Stock	Chrome	8.0
CRL															
Int. 1	Revoked	EV	✓	EV	✗	✓	✓	✓	✓	✓	✓	✗	✗	✗	✗
	Unavailable	EV	✓	—	✗	✗	✓	✓	✓	✓	✓	✗	✗	✗	✗
Int. 2+	Revoked	EV	EV	EV	✗	✓	✓	✓	✓	✓	✓	✗	✗	✗	✗
	Unavailable	✗	✗	—	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗
Leaf	Revoked	EV	EV	EV	✗	✓	✓	✓	✓	✓	✓	✗	✗	✗	✗
	Unavailable	✗	✗	—	✗	✗	✗	✗	✗	A	✓	✗	✗	✗	✗
OCSP															
Int. 1	Revoked	EV	EV	EV	EV	✗	✓	✓	✓	✓	✓	✗	✗	✗	✗
	Unavailable	✗	✗	—	✗	✗	L/W	✗	✓	✓	✓	✗	✗	✗	✗
Int. 2+	Revoked	EV	EV	EV	EV	✗	✓	✓	✓	✓	✓	✗	✗	✗	✗
	Unavailable	✗	✗	—	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗
Leaf	Revoked	EV	EV	EV	✓	✓	✓	✓	✓	✓	✓	✗	✗	✗	✗
	Unavailable	✗	✗	—	✗	✗	✗	✗	✗	A	✓	✗	✗	✗	✗
Reject unknown status		✗	✗	—	✓	✓	✗	✗	✗	✗	✗	—	—	—	—
Try CRL on failure		EV	EV	—	✗	✗	L/W	✓	✓	✓	✓	—	—	—	—
OCSP Stapling															
Request OCSP staple		✓	✓	✓	✓	✓	✓	✗	✓	✓	✓	✗	I	I	✗
Respect revoked staple		✗	✓	—	✓	✓	L/W	—	✓	✓	✓	—	—	—	—

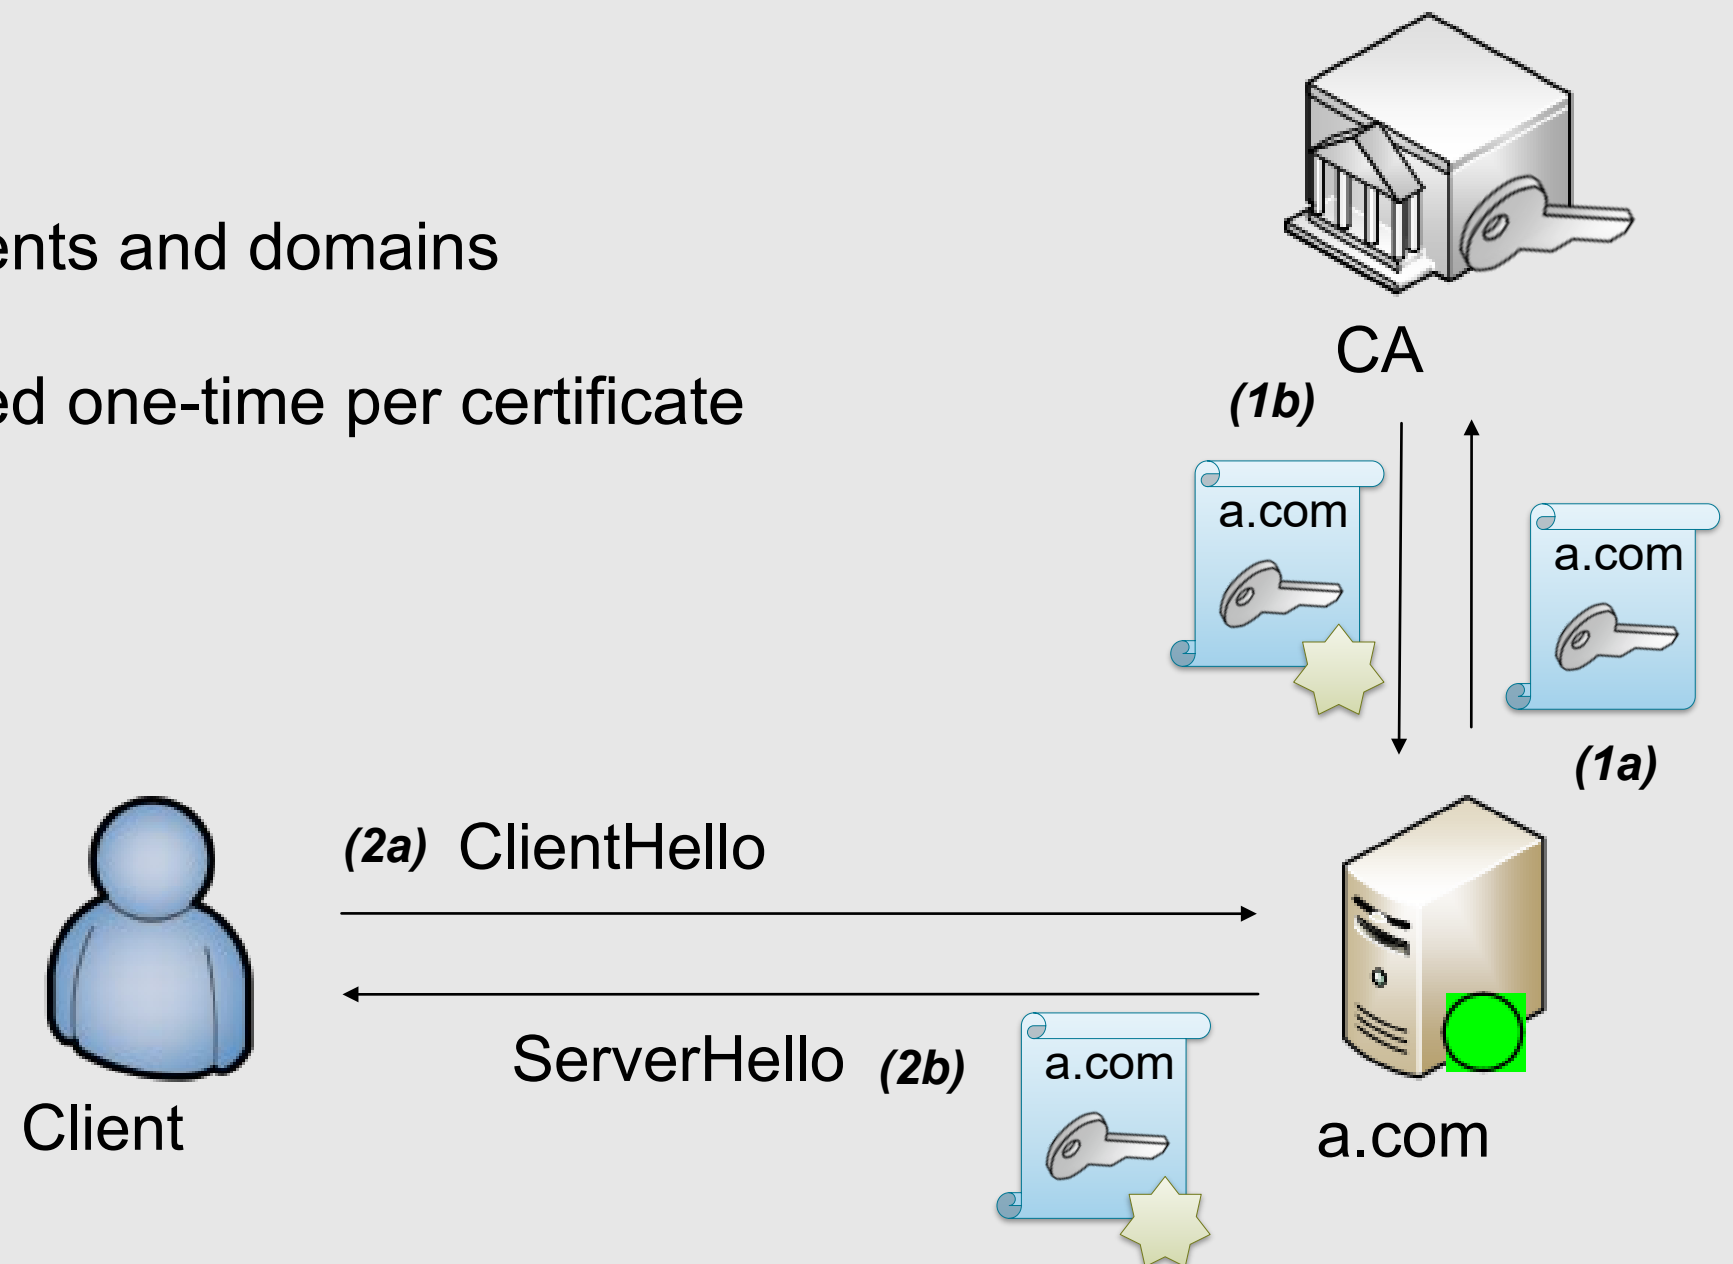
Table 2: Browser test results, when intermediate (Int.) and leaf certificates are either revoked or have revocation information unavailable. ✓ means browser passes test in all cases; ✗ means browser fails test in all cases. Other keys include EV (browser passes only for EV certificates), L/W (browser passes only on Linux and Windows), A (browser pops up an alert), and I (browser requests OCSP staple but ignores the response).

Certificate-Chain Validation

- The root CA certificate is trusted.
- All certificates are valid.
 - ✓ $\text{NotBefore} < \text{time}() < \text{NotAfter}$
 - ✓ Not revoked (if revocation is supported)
- The leaf certificate is issued for the contacted party.
- Certificates form a *chain of trust*.
 - ✓ 1st certificate is self signed, and i th certificate's issuer is $(i-1)$ th certificate's subject.
 - ✓ 2nd certificate can be verified with the public key of the 1st one, 3rd certificate can be verified with the public key of the 2nd one,..., i th certificate can be verified with the public key of the $(i-1)$ th.


SSL/TLS PKI Model

- SSL/TLS Protocol
- CA is trusted by clients and domains
- Step (1) is performed one-time per certificate



SSL/TLS PKI: Weak Authentication

- Certificates signed by single CA
 - ✓ Currently, cannot sign certificate by multiple CAs.
- Weakest-link security with too many *trusted* entities
 - ✓ Current browsers trust ~1500 keys that can issue valid certificates.

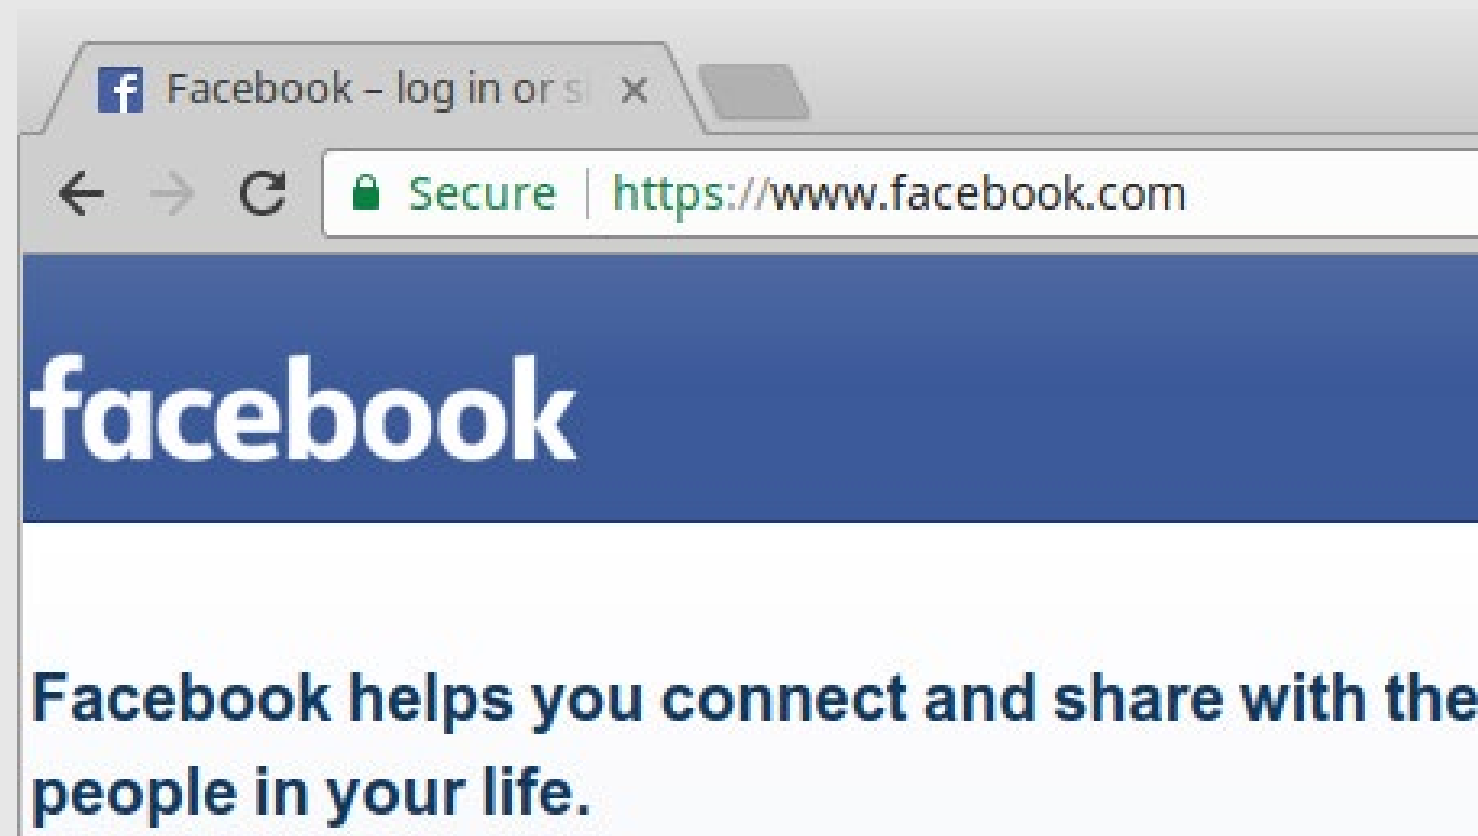


SSL/TLS PKI Problems

- Weakest-link security
- Revocation system is insecure and inefficient
 - ✓ Various schemes
 - ✓ Some CAs are *too-big-to-fail*
- Trust agility
 - ✓ Domains cannot state which CAs are trusted
- Transparency
 - ✓ CAs' actions are not transparent
- Imbalance
 - ✓ CAs have almost unlimited power
- Misconfigurations
 - ✓ SSLv2, weak crypto, NULL cipher suites

SSL/TLS as a Secure Channel

- Secure communication
 - ✓ Client-Server via HTTPS



PKI vs Key Server

- **Pros & Cons**

- ✓ Key server requires everyone online in real time. For PKI, CA need not be online in real time (unless for OCSP).
- ✓ Key server is a single point of failure. For PKI, CRL database is less security-critical and easier to distribute.
- ✓ Key server distributes symmetric keys which cannot be used for non-repudiation. With PKI, private key owner cannot deny its digital signatures.
- ✓ Key server needs the master key in online computer. For PKI, the root CA's private key is rarely used and need not in online computer.
- ✓ PKI is much more complex than key server, and requires more computational power.

Key Points

- PKI addresses key management problem.
- Trust models (Monopoly, Oligarchy, Anarchy) will decide how PKI is established.
- Certificate revocation is the hardest problem to solve in PKI.

Exercises & Reading

- Classwork (Exercise Sheet 12): due on Fri Nov 30, 10:00 PM
- Homework (Exercise Sheet 12): due on Fri Dec 7, 6:59 PM
- Reading: FSK [Ch18, Ch19, Ch20]
- Final exam (Week 14): Fri 14 Dec, 7:00-9:30 PM (mainly covering Week 5 – Week 12, open-book but no Internet access)

Week 13 Presentations

Slides are available on eDimension.

1. Liu Bowen

- ✓ **Title:** *Enter the Hydra: Towards Principled Bug Bounties and Exploit-Resistant Smart Contracts* (Usenix Security 2018)

2. Sakshi Sunil Udeshi

- ✓ **Title:** *Turning Your Weakness Into a Strength: Watermarking Deep Neural Networks by Backdooring* (Usenix Security 2018)

3. Flavio Toffalini

- ✓ **Title:** *The Guard's Dilemma: Efficient Code-Reuse Attacks Against Intel SGX* (Usenix Security 2018)

4. Tok Yee Ching

- ✓ **Title:** *Things You May Not Know About Android (Un)Packers: A Systematic Study based on Whole-System Emulation* (NDSS 2018)

Review of Weeks 5-12

Topics

- **Week 5: Symmetric Encryption**
 - ✓ Attacks (COA, KPA, CPA, CCA) & security level
 - ✓ Block ciphers (AES)
 - ✓ Block cipher modes (ECB, CBC, OFB, CTR)
- **Week 8: Hash & MAC**
 - ✓ Merkle-Damgard construction
 - ✓ MD-based hash functions
 - ✓ Hash-based MACs (HMAC)
 - ✓ Cipher-based MACs (CBC-MAC, CMAC)

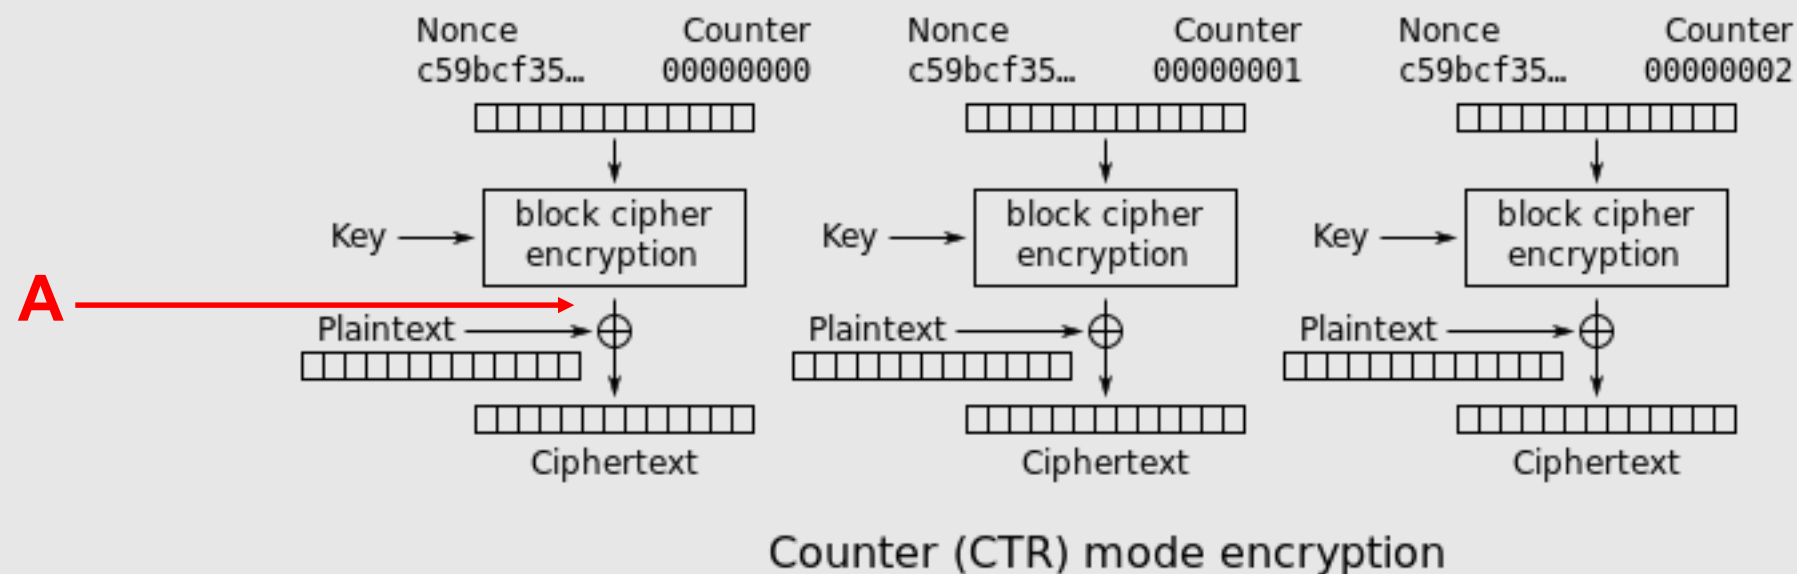
Topics

- **Week 9: Secure Channel & Randomness**
 - ✓ Order of authentication and encryption
 - ✓ Authenticated encryption (CCM, GCM)
 - ✓ Real randomness vs Pseudo-randomness
- **Week 10: Public-Key Cryptography**
 - ✓ Computations modulo a prime
 - ✓ (E)GCD algorithms
 - ✓ Cyclic groups & CRT
 - ✓ Diffie-Hellman algorithm
 - ✓ RSA algorithm

Topics

- **Week 11: Cryptographic Protocols**
 - ✓ Key negotiation (authenticated DH)
 - ✓ Key distribution (Kerberos)
 - ✓ Entity authentication (symmetric key based & public key based)
 - ✓ Non-repudiation (fairness, TTP involvement)
- **Week 12: PKI**
 - ✓ Trust models
 - ✓ Digital certificates & revocation

W5-HW6: An adversary observes the communication encrypted using CTR mode with **the same fixed nonce**. The nonce is hardcoded, so it is not included in the ciphertext. The adversary knows the 16-byte ciphertext **C** and **C'**, and the plaintext **P** corresponding to C. What information, if any, can the adversary infer about the plaintext **P'** (corresponding to C')?



- Because using the same nonce,

$$P \oplus A = C$$

$$P' \oplus A = C'$$
- Then $P \oplus P' = C \oplus C'$
- If the adversary knows C, C' and P are know, it can easily recover P'.

W8-HW4: Suppose message a is one block long. Suppose that an attacker has received the **MAC** t for a using CBC-MAC under some random key unknown to the attacker. Explain how to forge the MAC for a two-block message of your choice. What is the two-block message that you chose? What is the tag that you chose? Why is your chosen tag a valid tag for your two-block message?

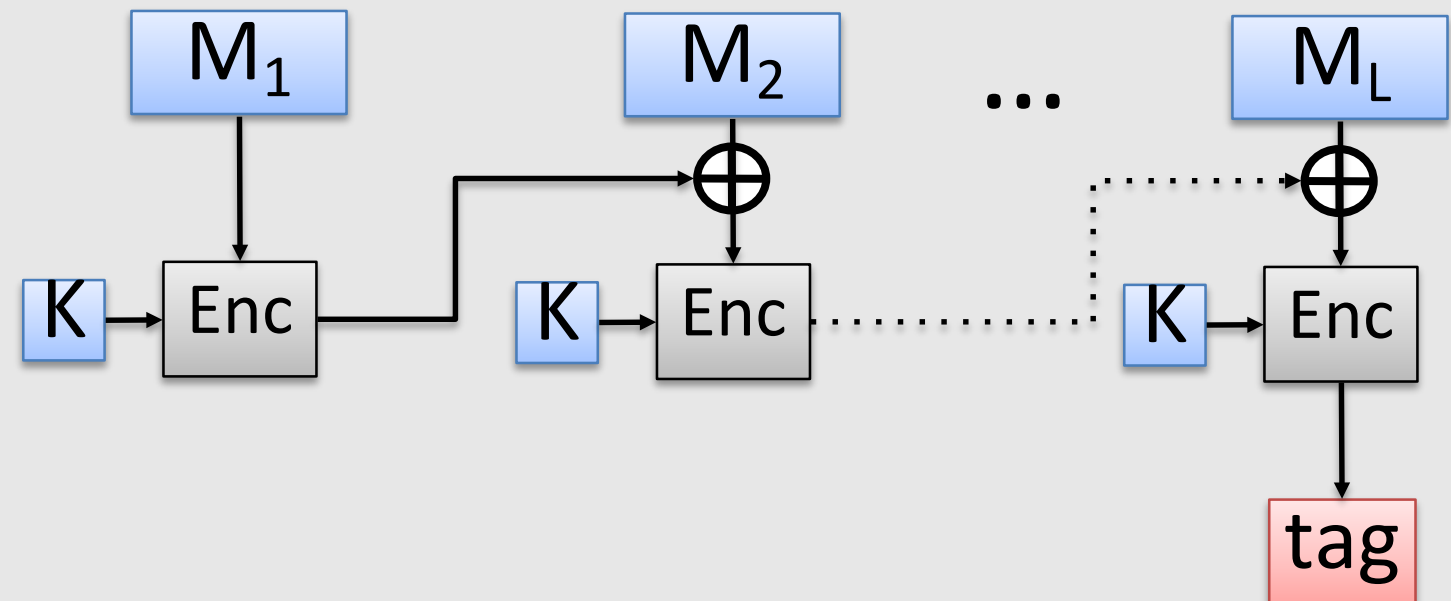
- As a is one-block message,
 $MAC_K(a) = E_K(a) = t$
- The two-block message to be chosen is:

$$M = a || (t \oplus a)$$

- The tag can be calculated as follows:

$$\begin{aligned} MAC_K(M) &= MAC_K(a || (t \oplus a)) \\ &= E_K((t \oplus a) \oplus t) \\ &= E_K(a) \\ &= MAC_K(a) \\ &= t \end{aligned}$$

- The chosen tag t is valid for the chosen message M as shown in the above steps.



$$C_1 = \text{Enc}(K, M_1)$$

$$C_2 = \text{Enc}(K, M_2 \oplus C_1)$$

$$C_3 = \text{Enc}(K, M_3 \oplus C_2)$$

...

$$C_L = \text{Enc}(K, M_L \oplus C_{L-1})$$

$$\text{tag} = C_L$$

W10-HW1: Proof of $\text{lcm}(a,b) = ab/\text{gcd}(a,b)$

- By taking the prime power decomposition of a and b , we have

$$a = p_1^{c_1} \times p_2^{c_2} \times \dots \times p_k^{c_k}$$

$$b = p_1^{d_1} \times p_2^{d_2} \times \dots \times p_k^{d_k}$$

where each of the p_i are distinct primes and each of the c_j and d_j are non-negative integers.

- The important part of this trick is that we write both a and b as a product of the same primes, even if some of the powers are zero. By definition:

$$\text{lcm}(a,b) = p_1^{\max(c_1,d_1)} \times \dots \times p_k^{\max(c_k,d_k)}$$

$$\text{gcd}(a,b) = p_1^{\min(c_1,d_1)} \times \dots \times p_k^{\min(c_k,d_k)}$$

- We notice that $\max(c_i,d_i) + \min(c_i,d_i) = c_i + d_i$

$$\text{lcm}(a,b) \times \text{gcd}(a,b) = p_1^{c_1+d_1} \times \dots \times p_k^{c_k+d_k}$$

$$= (p_1^{c_1} \times p_1^{d_1}) \times \dots \times (p_k^{c_k} \times p_k^{d_k})$$

$$= a b$$

- Therefore $\text{lcm}(a,b) = ab/\text{gcd}(a,b)$

W11-HW2: Fair Non-repudiation Protocol Using Off-line TTP

- **IEEE CSFW'97** (simplified)

1. $A \rightarrow B$: C, EOO_C

2. $B \rightarrow A$: EOR_C

3. $A \rightarrow B$: K, EOO_K

IF 4. $B \rightarrow A$: EOR_K **THEN stop**

ELSE {3'. $A \rightarrow TTP$: K

4'. $B \leftarrow TTP$: K, con_K

5'. $A \leftarrow TTP$: con_K }

- **Evidence**

✓ B receives EEO_C, EEO_K if TTP is not involved, otherwise EEO_C, con_K

✓ A receives EOR_C, EOR_K if TTP is not involved, otherwise EOR_C, con_K

- K – message key defined by A

- $C = Enc_K(M)$ – cipher text of M

- EEO_C, EOR_C – evidence of origin and receipt of C

- EEO_K, EOR_K – evidence of origin and receipt of K

- con_K – evidence of confirmation of K

- **Problem:** B may not be able to terminate the protocol run timely without breaching fairness.

- Further improve ?

End of Slides for Week 12